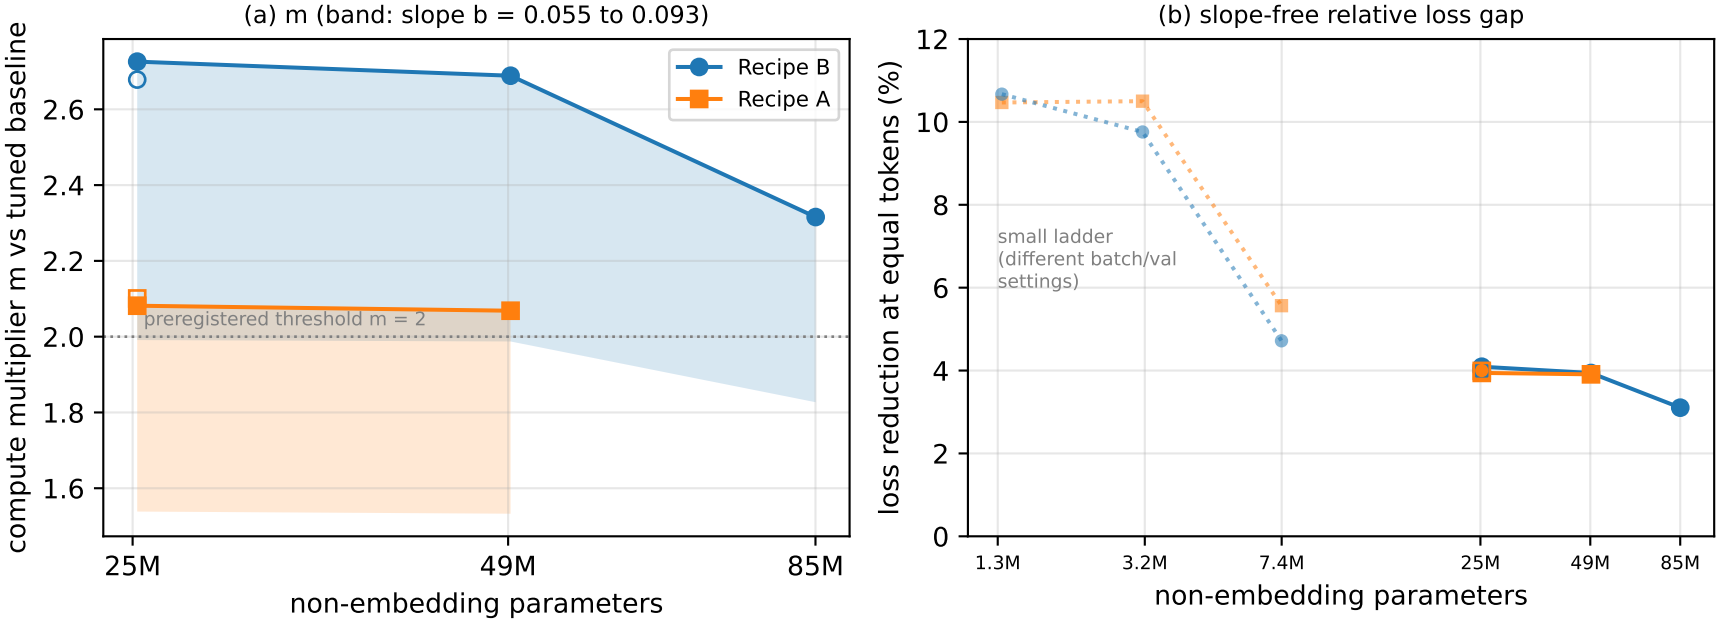
(a) m (band: slope b = 0.055 to 0.093)
compute multiplier m vs tuned baseline
Recipe B
Recipe A
preregistered threshold m = 2
2.6
2.4
2.2
2.0
1.8
1.6
25M
49M
85M
non-embedding parameters
(b) slope-free relative loss gap
loss reduction at equal tokens (%)
small ladder
(different batch/val
settings)
12
10
8
6
4
2
0
1.3M
3.2M
7.4M
25M
49M
85M
non-embedding parameters